

The artifact

- **Repository:** <https://github.com/yitaohou/reign-watch>
- **Walkthrough (4 min):** <https://www.loom.com/share/cfc7f9603f80419fa43589335b37c07a>

Reign Watch — a Claude agent, a custom MCP server, and HubSpot, wired so that a change in a public regulation becomes one sourced, account-specific brief per affected bank, with an R-17 audit record written before anything else and a named human between the brief and any send. Code and setup are in `artifact/` and `README.md`.

FLOW

```
scheduler tick / "Check now"
|
v  claude -p · skill: regulatory-watch
|  tools: reign-tools MCP + HubSpot MCP (read-only)
|
+-- checking    diff_source() per URL
|               fetch -> extract text or links -> hash -> compare to last snapshot
|               -> changed?  added/removed lines · new links · first_seen
|
+-- judging     the agent reads the diff
|               all first_seen ..... baseline      snapshot stored, nothing drafted
|               nothing changed ..... no_change
|               not material ..... noise           reason + raw diff kept; no action
|               material ..... continue
|
+-- screening   companies from HubSpot, each one in this order:
|               CEO exclusion -> segment not reached -> fails an ICP gate
|               -> fails the regulation's flag -> otherwise AFFECTED
|               one reason per company; a `score` record per affected FS account
|
+-- writing      emit_brief() once per affected account
|               write_audit (create)  fails closed: no record, no brief
|               render                 section 1 shared · section 2 from that
|                                     account's record · one ask
|               brief_eval (7 checks) citations · <=250 words · banned phrases
|                                     · source hashes · audit precedes brief
|                                     · send=false · one ask
|
|               handoff.json          only on PASS
|
|
v
dashboard      live phase while running -> result when done
               -> briefs wait for the named approver
```

Who does what

PART	RESPONSIBILITY
Server (<code>server/app.py</code>)	Decides <i>when</i> a check runs and for which regulation; starts the agent as a headless <code>claude -p</code> ; reads its event stream to show the live phase; records each cycle; serves the dashboard and the approval / kill API. Makes no judgment.

PART	RESPONSIBILITY
Tools (<code>mcp/reign_tools.py</code>)	The only code that touches the outside or the disk: fetch a page, diff it against the last snapshot, look up a person or account, write an R-17 record, evaluate kill criteria, turn a brief into a file. <code>write_audit</code> refuses incomplete records; <code>emit_brief</code> cannot render before the record exists.
Watch skill (<code>skill/regulatory-watch</code>)	The agent's procedure for one cycle: which URLs to diff, how to classify the cycle, the materiality test, the screening order, when to call the brief skill, the JSON it must end with.
Brief skill (<code>skill/regulatory-brief</code>)	The writing rules: three sections, every sentence cited to a fetched source, the account's CRM record only and phrased as a condition, one ask, no banned phrases, 250 words.
Eval (<code>eval/brief_eval.py</code>)	Seven mechanical checks run inside <code>emit_brief</code> ; a failing brief is renamed <code>.FAILED.html</code> and gets no hand-off.
HubSpot	System of record. Companies and their <code>reign_*</code> flags are read during screening; the agent never writes to it.
Dashboard (<code>server/dashboard.html</code>)	Regulations → Checks → one Check → Brief; Audit page; who you are acting as; Approve (named approver only); Kill (CEO / CRO only); test mode.
Named humans (<code>data/directory.json</code>)	Approve a brief; kill the motion. Nothing else in the system can do either.

Python fetches, diffs, records and blocks. The agent reads, judges, screens and writes. A person approves or kills.

WHERE THE CONSTRAINTS LIVE

- **R-17** — `write_audit` refuses any record missing a field, with a vague purpose, or with an unresolvable principal, approver or object; `@audited` makes the brief's render code unreachable until the record is written; eval check 5 verifies independently that the record predates the brief.
- **Named human** — `approver` must resolve to a person in the directory; the dashboard enables Approve only for that person and the API refuses anyone else; Kill is visible only to `can_kill` humans.
- **Precision** — one change fans out to N accounts as N distinct briefs, each written from that account's own record and each approved separately; every account screened out is listed with its reason.
- **Reusability** — one `playbook_id` ; a new regulation is a new trigger file and a playbook instance; a new product launch is the same playbook with `product` and `audience` changed and `version` bumped.

TEST MODE

The watched documents rarely change, so the dashboard offers two scenarios that seed index-page baselines from Wayback Machine captures (January, April, August, September 2026) in a separate state tree and let a check replay what actually happened: scenario A surfaces SR 26-2 itself on the Fed index and produces a brief per affected account; scenario B surfaces SR 26-6 and is judged noise. Normal-mode state is untouched.